

# Results of some performance and security evaluation against library-based compartments created in the Morello Board

Regis Schuch<sup>1</sup>, Rafael Z. Frantz<sup>1</sup>, Carlos Molina-Jiménez<sup>3</sup>

<sup>1</sup>Unijuí University – Ijuí – Brazil

<sup>2</sup>University of Cambridge – Cambridge – United Kingdom

{regis.schuch, rzfrantz}@unijui.edu.br,

carlos.molina@cl.cam.ac.uk

**Abstract.** *This document presents the results of some performance and security evaluation tests conducted on library-based compartments created on a Morello Board running cheriBSD 24.5.*

## 1. Evaluation of the number of library-based compartments

In the CAMB project we are interested in using Morello Boards for creating attestables (compartments) to support the execution of exfiltration sensitive processes. The idea is to create attestables to host the executions of the sensitive code of the exfiltration sensitive process. We account for attestables that remain active for several hours or days. The question that emerges here is how many compartments can be created to run simultaneously on a single Morello Board? To find this limit, we created and run simultaneously on a Morello Board an increasing number of compartments. Within each compartment we deployed and integration process that sent requests to remote applications to retrieve data. We thought that the limit depended on the amount of RAM memory of the Morello Board and tried to exhaust it by means of systematically increasing the number of compartments.

### 1.1. Experiment

The main aim of this experiment is to measure and analyse how the memory of a Morello Board is consumed by instances (also called replicas) of attestables. To this end, we loaded the attestable with a C program compiled with the library compartmentalisation tool and loaded in a compartment; precisely, we used the enterprise application integration (see yellow box) use case implemented in - tee-compartmentalisation-study-case repository.

The metric to measure is the number of attestables that can be created on a Morello Board before consuming 90% of its memory.

In addition to the number of attestables, we took the opportunity to collect metrics about the time it takes the operating system to wipe the memory used by the attestable.

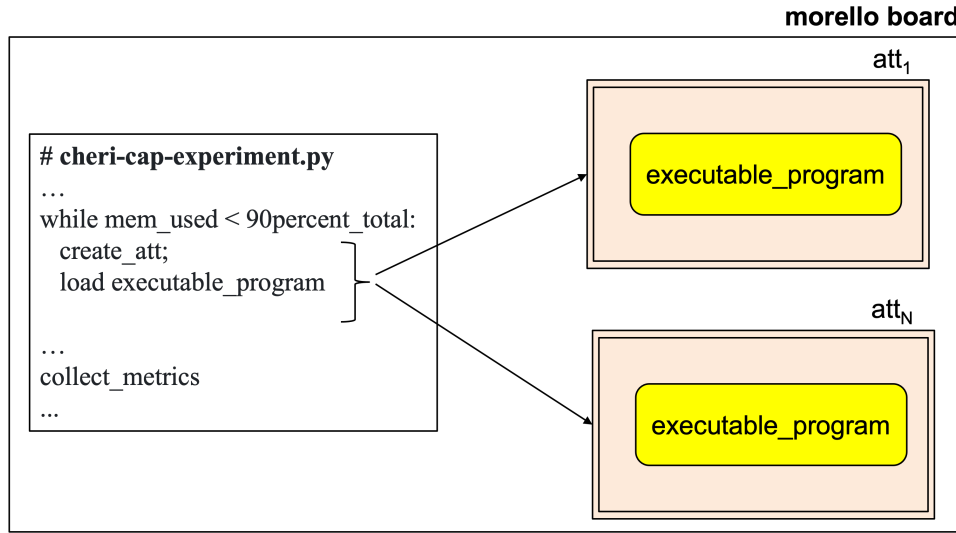
Some experimental facts:

- (1) The Morello board used to conduct the experiment has 17,118,408,704 bytes (approximately 17,118.4 MB). Thus, 90 per cent of its memory is 15,406,567,833.6 bytes (approximately 15,406.5 MB).
- (2) In the experiments that we conducted, we loaded code of the EAI implemented in (see yellow box) - tee-compartmentalisation-study-case repository. We compiled as shown below:

```
$ clang-morello -march=morello+c64 -mabi=purecap -g -o
integration_process integration_process.c -L. -
lssl -lcrypto -lpthread
```


- (3) `cheri-cap-experiment.py` script is used to create the replicas of the attestables, and collect metrics. We incremented the number of replicas created from 1 to N. See replication of attestable results.

The Figure 1 shows the experiment set up:



**Figure 1. Memory exhaustion by attestable replicas.**

The graph, shown in Figure 2, shows the behaviour of memory consumption and elapsed time as the number of testable replicas increases on the Morello Board.



**Figure 2. Memory consumption and elapsed time versus number of replicas on the Morello board.**

Imagine that user Alice is conducting the experiment. To create the attestables and collect the metrics, Alice executes the following steps:

(1) **Initiation:** In the Morello board, Alice initiates *cheri-cap-experiment.py*.

(2) **Launch:** To launch *cheri-cap-experiment.py*, Alice executes:

```
% python3 cheri-cap-experiment.py
```

(3) `python3 cheri-cap-experiment.py` runs incrementally, creating attestable replicas until it detects that the attestables have consumed 90% of the 17,118.4 MB of the Morello Board's memory, that is, about 15,406.5 MB.

### **Preliminary observations:**

The results are shown in replication of attestable results and exhibit an unexpected behaviour. We expected memory consumption to increase steadily from 1,628.3 MB, which corresponds to a single attestable replica, to 15,406.5 MB (90% of total memory) consumed by N attestable replicas. The aim was to determine the value of N.

However, unexpectedly, memory consumption increased steadily only up to the creation of 3,800 attestable replicas, which consumed 14,582.5 MB. Beyond that, intriguingly, memory consumption decreases as the number of attestable replicas increases. Note that the last metric shows that 8,991 attestable replicas consume 13,066.4 MB (approximately 76% of the total memory).

We do not have an explanation for that yet. We need to double check measurements. We are showing this preliminary result only to show that this is a pending question in our agenda. The analysis of the time to wipe the memory of the attestable replicas is also pending.

The `cpu-performance` and `memory-performance` folders contain code and results of some experiments in progress. `Cpu-performance` is meant to measure the performance of code running inside attestables in the computation of some mathematical operations, integer arithmetic, floating point arithmetic, and array manipulation. `Memory-performance` is meant to assess the time it takes an attestable to allocate and free blocks of memory of different sizes.

## **2. CPU Performance Tests on the Morello Board**

The main aim of this experiment is to measure and analyse the performance of the Morello Board's CPU by comparing the results of identical tests carried out inside and outside a secure enclosure. The tests measured the time required to perform computational operations, including complex mathematical functions, arithmetic operations with integers, floating point operations and matrix manipulation. The complex mathematical functions included sine, cosine, tangent, square root and logarithm. The integer arithmetic tests focused on multiplication, division, subtraction and modulo operations. For floating point operations, the focus was also on multiplication, division and subtraction. The matrix manipulation test involved initialising matrices and performing operations such as multiplication and division. The execution time of each operation was measured and recorded. These tests were repeated 100 times for each operation, in both environments - inside the secure compartment and in the Morello Board's normal operating environment - and the results were saved in a CSV format file for each environment.

Algorithm 1 details how to carry out the CPU performance tests and record the results.

Execution begins with the `perform_tests` function (line 1), which receives as a parameter a log file where the results will be stored and the total accumulated time needed to run the tests. The function enters a repeat loop that is repeated the number of times specified by `NUM.TESTS` (line 3), where each iteration represents a test identified by `test_num`. In each iteration, the initial test time is recorded (line 4), followed by the execution of the computational

---

**Algorithm 1** CPUPerformance

---

```
1: perform_tests(log_file, total_time)
2: begin
3: for test_num in NUM_TESTS do
4:   start_time = capture_time()
5:   execute_operations(WORKLOAD_SIZE)
6:   end_time = capture_time()
7:   cpu_time = calculate_cpu_time(start_time, end_time)
8:   results(log_file, test_num, cpu_time)
9:   total_time += cpu_time
10: end for
11: end
```

---

operations determined by WORKLOAD.SIZE (line 5). At the end of execution, the final time is recorded (line 6), and the total CPU time used is calculated by subtracting the start\_time from the end\_time (line 7). This time is then recorded in the log file, along with the test number (line 8), and also added to total\_time, which accumulates the total time spent on all the tests (line 9).

The code used for the tests is available from the tee-morello-performance-experiments repository.

### Test procedure

Suppose user Alice is conducting the experiment. To carry out the CPU performance tests and collect the results, Alice takes the following steps:

1. **Start:** Alice compiles and runs the test program in the two scenarios described above:
  - **Inside the compartment:** Alice runs the program `cpu-in-experiment.c` on the Morello Board, using the secure environment.
    - **Compile:** `clang-morello -march=morello+c64 -mabi=purecap -g -o cpu-in-experiment cpu-in-experiment.c -L. -lm`
    - **Run:** `procontrol -m cheriC18n -s enable ./cpu-in-experiment`
  - **Outside the compartment:** Alice runs the program `cpu-out-experiment.c` in the Morello Board's normal operating environment.
    - **Compile:** `clang-morello -o cpu-out-experiment cpu-out-experiment.c -lm`
    - **Run:** `./cpu-out-experiment`
2. **Execution:** The program iterates automatically through the different types of operations, performing complex mathematical calculations, arithmetic operations with integers, floating point operations, and matrix manipulation.
3. **Repetition:** Each operation is repeated 30 times. The time of each operation is recorded in files in CSV format for both environments.

### Results

The results of the tests carried out inside the secure compartment were stored in the file `cpu-in-experiment-result.csv`, while the results of the run in the Morello Board's

normal environment were stored in `cpu-out-experiment-result.csv`. The Table 1 compares the average execution times for each type of operation in the two environments.

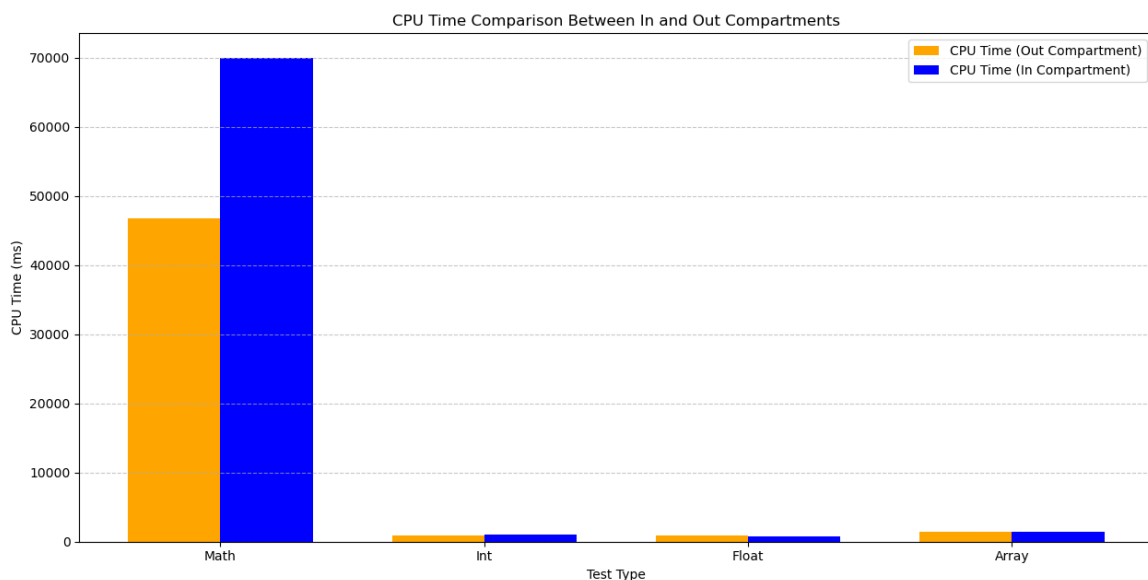
**Table 1. CPU Time Comparison**

Test Type	CPU Time (ms) - Out Compartment	CPU Time (ms) - In Compartment
Maths	46,696	69,998
Int	923	993
Float	816	785
Array	1,419	1,460

### Analysing the results

According to the Table 1, the results show that, on average, complex mathematical operations within the secure compartment took 69,998 ms, while in the normal operating environment the time was 46,696 ms, representing a difference of approximately 49.74 per cent. For arithmetic operations with integers, the average time was 993 ms inside the secure compartment and 923 ms in the normal operating environment, a difference of 7.58 per cent. For floating point operations, the time inside the safe compartment was 785 ms compared to the normal operating environment (816 ms), showing a slight improvement of 3.80 per cent. As for matrix manipulation operations, the time inside the safe compartment was 1,460 ms, while in the normal operating environment it was 1,419 ms, indicating an increase of 2.89 per cent.

The Figure 3 illustrates the differences in performance between operations conducted within and outside the secure compartment.



**Figure 3. Comparison of CPU performance times in and out of the secure compartment on the Morello board.**

These results suggest that the Morello Board's safe compartment introduces a small performance overhead in more complex operations, such as mathematical calculations. However, for simpler operations, such as floating point, the difference is minimal, with a slight increase in execution time within the secure compartment. This indicates that the safe compartment can be

used efficiently in specific scenarios without significantly compromising performance, although the impact is more noticeable in operations that require more processing.

### 3. Memory Performance Tests on the Morello Board

The main aim of this experiment is to measure and analyse the performance of memory operations on a Morello Board by testing memory blocks of different sizes. Specifically, the test evaluates the time required to allocate, write, read and free memory blocks. To carry out this experiment, a C program was implemented - according to the formalism in Algorithm 2 - which automates the memory testing process. The programme runs repeated tests, increasing the size of the blocks at each stage and recording the times of each operation. The results are stored in a file in CSV format. The tests were carried out in two different environments: inside the TEE, where the programme was compiled for the CHERI architecture and run in a secure compartment of the Morello Board; and outside the TEE, where the programme was compiled and run in the normal operating environment of the Morello Board.

---

#### Algorithm 2 MemoryPerformance

---

```

1: perform_tests(log_file, total_time)
2: begin
3:   for block_size in MIN_BLOCK_SIZE to MAX_BLOCK_SIZE step BLOCK_STEP do
4:     for test_num from 1 to NUM_TESTS do
5:       allocation_time = time(malloc(block_size))
6:       write_time = time(write_to_memory(block, block_size))
7:       read_time = time(read_from_memory(block, block_size))
8:       free_time = time(free(block))
9:       log(log_file, block_size, test_num, allocation_time, write_time, read_time, free_time)
10:    end for
11:  end for
12: end

```

---

Algorithm 2 details how to perform memory performance tests and record the results. Execution begins with the `perform_tests` function (line 1), which receives as parameters a log file where the results will be stored and the total accumulated time needed to run the tests. The function enters a repeat loop that iterates over different memory block sizes, defined from `MIN_BLOCK_SIZE` to `MAX_BLOCK_SIZE` with increments specified by `BLOCK_STEP` (line 3). For each block size, another loop is started, repeating the tests the number of times defined by `NUM_TESTS` (line 4). At each iteration, the memory allocation time is measured with the `malloc` function (line 5), followed by the time to write to the block (line 6), the time to read the block (line 7) and, finally, the time to free the memory with the `free` function (line 8). These times are recorded in the log file along with the test number (line 9). After all the iterations for all the memory blocks, the process is completed, allowing the performance of the memory allocation, write, read and free operations to be measured and compared in different block size configurations.

### Experimental Configuration

1. The memory blocks tested range from 100 MB to 1 GB, with increments of 100 MB at each step. Each test involves 30 iterations per block size. The diagram in the Figure 4 shows the block pattern and the times for each operation associated with different block sizes.

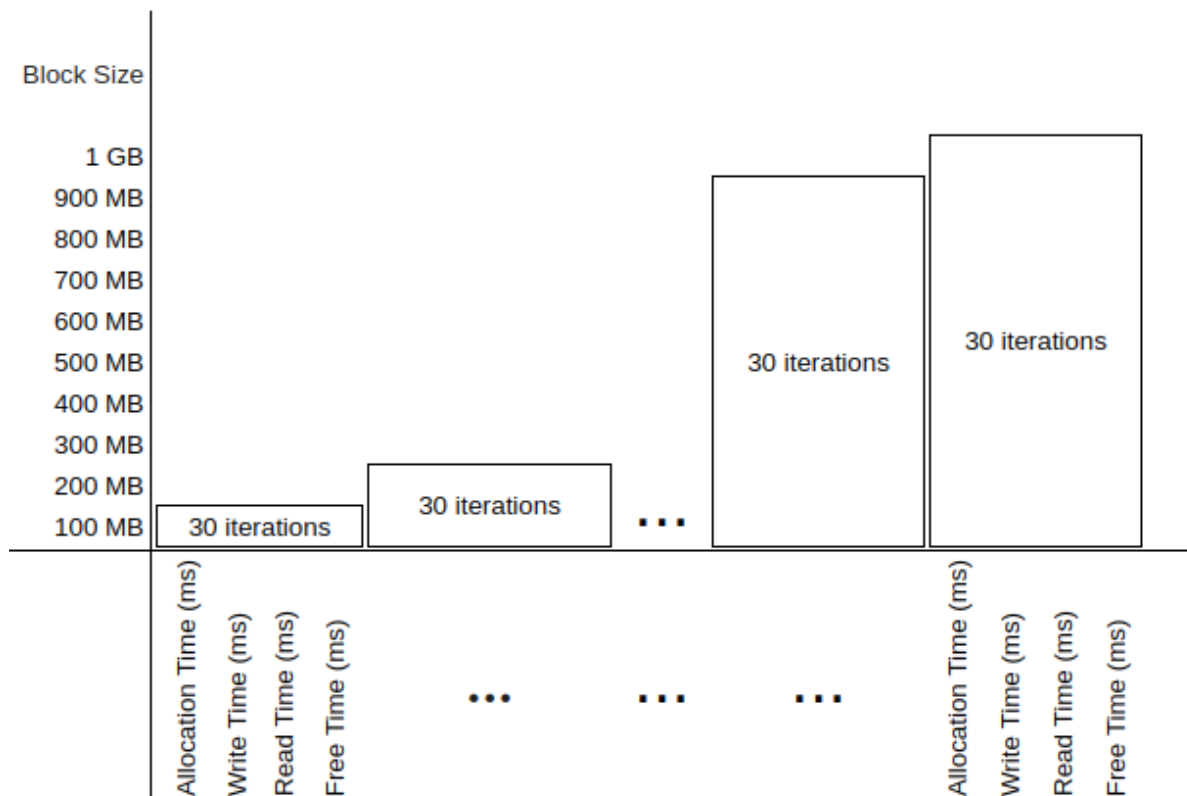


Figure 4. Memory performance at different block sizes on the Morello board.

The operations measured are:

- Memory allocation:** time required to allocate a block of memory.
  - Writing to Memory:** time required to write data to the entire allocated memory block.
  - Memory read:** time taken to read the data from the entire memory block.
  - Memory release:** time taken to release the memory block back into the system.
- The code for the experiment we carried out is available in the tee-morello-performance-experiments repository.

## Test procedure

Imagine that user Alice is conducting the experiment. To carry out the memory performance tests and collect the results, Alice performs the following steps:

- Start:** Alice compiles and runs the test programme in two different scenarios:
  - Inside the compartment:** Alice runs the programme `memory-in-experiment-result.c` on the Morello Board, using the secure environment.
    - Compile:** `clang-morello -march=morello+c64 -mabi=purecap -g -o memory-in-experiment memory-in-experiment.c -L. -lm`
    - Run:** `procontrol -m chericl8n -s enable ./memory-in-experiment`
  - Outside the compartment:** Alice runs the programme `memory-out-experiment-result.c` in the normal operating environment of the Morello Board.

- **Compile:** `clang-morello -g -o memory-out-experiment memory-out-experiment.c -lm`
  - **Run:** `./memory-out-experiment`
2. **Execution:** The programme automatically iterates over the memory block sizes, from 100 MB to 1 GB, performing the following sequence for each block size: allocate the block, write to the block, read the block, and release the block.
  3. **Repetition:** For each block size, these steps are repeated 30 times, both inside and outside the compartment. The time of each operation is recorded and saved in a CSV file for the environment.

## Results

For each CSV file, the iteration averages were calculated for each memory block size in the Allocation Time, Write Time, Read Time and Free Time attributes. For each memory block value (100 MB, 200 MB, etc.), the times measured in the iterations were added together and the average was obtained for each operation (allocation, write, read and free). This was done separately for the two files, representing in-compartment and out-of-compartment executions, respectively, as shown in the Tables 2 and 3.

**Table 2. In Compartment**

Block Size (MB)	Allocation Time (ms)	Write Time (ms)	Read Time (ms)	Free Time (ms)
100	106	295.308	282.576	97
200	138	590.498	565.152	272
300	138	885.784	847.719	318
400	141	1,180.815	1,130.297	492
500	131	1,476.014	1,412.881	321
600	168	1,771.086	1,695.456	399
700	251	2,066.147	1,978.022	725
800	235	2,361.646	2,260.586	738
900	312	2,656.590	2,543.166	1,197
1000	265	2,951.487	2,825.741	405

## Analysing the results

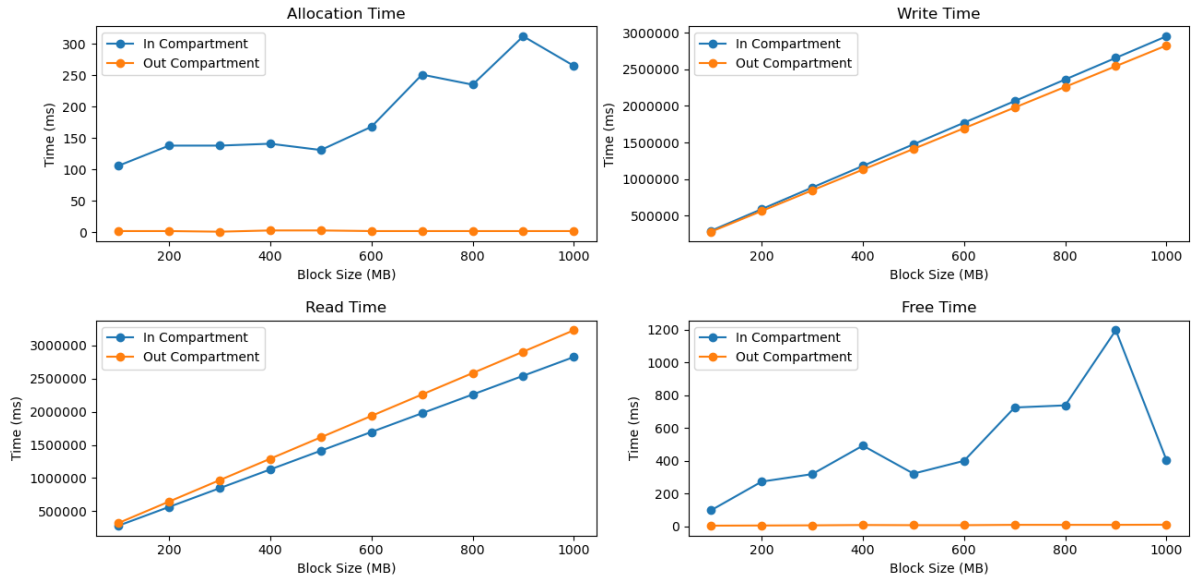
The results illustrated in the graphs in Figure 5 reveal some trends in relation to the four main operations: allocating, writing, reading and freeing memory.

- **Allocation time:** the values resulting from the tests in the secure compartment range from 106 ms for 100 MB blocks to 251 ms for 700 MB blocks, with some variations across block sizes. In the Morello Board’s normal operating environment, the allocation time is considerably shorter, ranging from 5 ms (100 MB) to 7 ms (900 MB).
- **Write time:** in both environments, the times follow a linear behaviour as the block size increases. In the secure compartment, the values start at 295,308 ms for 100 MB blocks and reach 2,951,487 ms for 1 GB blocks. In the normal operating environment, the times are slightly longer, ranging from 491,512 ms (100 MB) to 4,903,282 ms (1 GB).
- **Read time:** the results also show a linear behaviour in both environments. Inside the secure compartment, the read time starts at 282,576 ms (100 MB) and increases up to



**Table 3. Out Compartment**

Block Size (MB)	Size	Allocation Time (ms)	Write Time (ms)	Read Time (ms)	Free Time (ms)
100		2	282.574	322.944	3
200		2	565.140	645.880	4
300		1	847.708	968.814	5
400		3	1,130.294	1,291.766	7
500		3	1,412.856	1,614.708	6
600		2	1,695.426	1,937.630	6
700		2	1,977.999	2,260.573	8
800		2	2,260.593	2,583.527	8
900		2	2,543.151	2,906.485	8
1000		2	2,825.742	3,229.433	9

**Figure 5. Performance analysis of memory operations on the Morello board.**

2,825,741 ms (1 GB). In the normal operating environment, the times vary from 245,263 ms (100 MB) to 2,452,597 ms (1 GB), being slightly longer.

- **Memory release time:** shows similar behaviour in both environments. In the secure compartment, the time varies from 97 ms (100 MB) to 1,197 ms (900 MB), while in the normal operating environment the times are lower, ranging from 4,499 ms (100 MB) to 44,750 ms (1 GB).

Therefore, the results show that the secure environment imposes a greater burden on allocation time, with significant variations compared to the normal operating environment. On the other hand, write time and read time follow similar patterns, with the secure environment showing slightly higher results. However, the memory release time stands out, being significantly higher in the secure compartment. These results indicate that although the secure compartment imposes some overhead, it can still be used efficiently for memory operations.

#### 4. Morello Board Pipe Communication Performance Tests

The main objective of this experiment is to measure the performance of communication between processes on the Morello Board using pipes, evaluating the time taken to write and read messages. The experiment was run both inside a secure compartment and outside the compartment, in the normal operating environment of the Morello Board. The aim is to compare the performance of these two scenarios and understand the impact of secure memory compartmentalisation on the communication time between processes.

To conduct this experiment, we developed a C programme that automates the exchange of messages between parent and child processes via a pipe, used here as a communication mechanism between processes. During the experiment, the child process is tasked with writing messages of a predetermined length to the pipe. At the same time, the parent process reads these messages, recording the write and read times. The programme has been specifically configured to test messages of a defined size, with each test timing the writing and reading times. The results of these tests are recorded in a file in CSV format. Algorithm 3 gives a detailed description of the experimental flow.

---

**Algorithm 3** Pipe Communication Performance

---

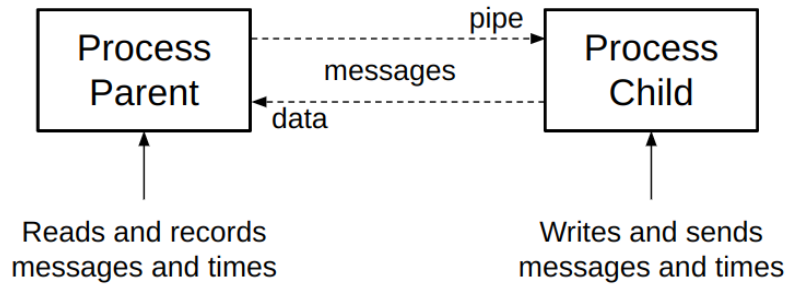
```
1: start_test(log_file)
2: begin
3: define STRLEN
4: define NUM_OF_MSG
5: for test_num from 1 to NUM_OF_MSG do
6:   if child_process then
7:     start_timer(write_time)
8:     write(pipe, message of size STRLEN)
9:     stop_timer(write_time)
10:    write(pipe, write_time)
11:   else
12:     read(pipe, message of size STRLEN)
13:     read(pipe, write_time)
14:     start_timer(read_time)
15:     stop_timer(read_time)
16:     log(log_file, test_num, write_time, read_time)
17:   end if
18: end for
19: end
```

---

In Algorithm 3, the start\_test function (line 1) starts a sequence of operations involving writing and reading messages through a pipe. First, the STRLEN and NUM\_OF\_MSG parameters are set in lines 3 and 4, establishing the size of the messages and the total number of messages to be sent. For each test, which iterates from 1 to NUM\_OF\_MSG (line 5), the child process, if active, starts timing the write time (line 7), writes a message of size STRLEN to the pipe (line 8) and then records the write time (line 9), sending this time back to the parent process via the pipe (line 10). In parallel, the parent process reads the message and the write time from the pipe (lines 12 and 13), starts counting the read time as soon as it starts reading (line 14), and stops counting when it finishes reading (line 15). The write and read times, along with the test number, are recorded in the log file (line 16). This process is repeated until all the tests have been completed (line 17).

## Experimental Configuration

1. The following diagram (Figure 6) represents the experimental configuration used to evaluate communication between processes via a pipe. The data blocks, in this case represented by messages, vary by a standard size of 1024 bytes per message, and a total of 100 messages are tested to analyse the consistency and efficiency of communication between internal processes.



### Configurations:

- Message Size (STRLEN): 1024 bytes
- Number of Messages (NUM\_OF\_MSG): 100

**Figure 6. Communication of the process via pipe.**

The operations measured are:

- a) **Writing to memory:** time taken to write data to the pipe.
  - b) **Memory read:** time taken to read the data sent via the pipe.
2. The code for the experiment we carried out is available in the tee-morello-performance-experiments repository.

## Test procedure

Imagine that user Alice is conducting the experiment. To carry out the pipe communication tests and collect the results, Alice performs the following steps:

1. **Start:** Alice compiles and runs the test programme in two different scenarios:
  - **Inside the compartment:** Alice runs the programme `pipe-in-experiment.c` on the Morello Board, using the secure environment.
    - **Compile:** `clang-morello -march=morello+c64 -mabi=purecap -g -o pipe-in-experiment pipe-in-experiment.c -L.`
    - **Run:** `proccontrol -m chericl8n -s enable ./pipe-in-experiment`
  - **Outside the compartment:** Alice runs the programme `pipe-out-experiment.c` in the normal operating environment of the Morello Board.
    - **Compile:** `clang-morello -g -o pipe-out-experiment pipe-out-experiment.c`
    - **Run:** `./pipe-out-experiment`

2. **Execution:** The programme sends a total of 100 messages, each of a predefined size of 1024 bytes. For each message sent by the child process, the write time is measured and the message, along with the write time, is transmitted via the pipe to the parent process, which in turn records the read time as soon as the message is received.
3. **Repetition:** Each set of write and read operations is repeated 100 times for each environment - inside the compartment and outside the compartment. The times of each operation are recorded in a file in CSV format for each environment.

## Results

The results of the tests carried out inside the secure compartment were stored in the file `pipe-in-experiment-result.csv`, while the results of the normal ambient execution of the Morello Board were stored in `pipe-out-experiment-result.csv`. The Tables 4 and 5 shows the records for each data set:

**Table 4. Performance Data Inside the Compartment**

Test	Message Size (Bytes)	Write Time (ms)	Read Time (ms)	Total Time (ms)
1	1024	0.016	0.161	0.177
2	1024	0.003	0.068	0.071
3	1024	0.003	0.075	0.078
4	1024	0.003	0.077	0.080
...	...	...	...	...
100	1024	0.003	0.079	0.082

**Table 5. Performance Data Outside the Compartment**

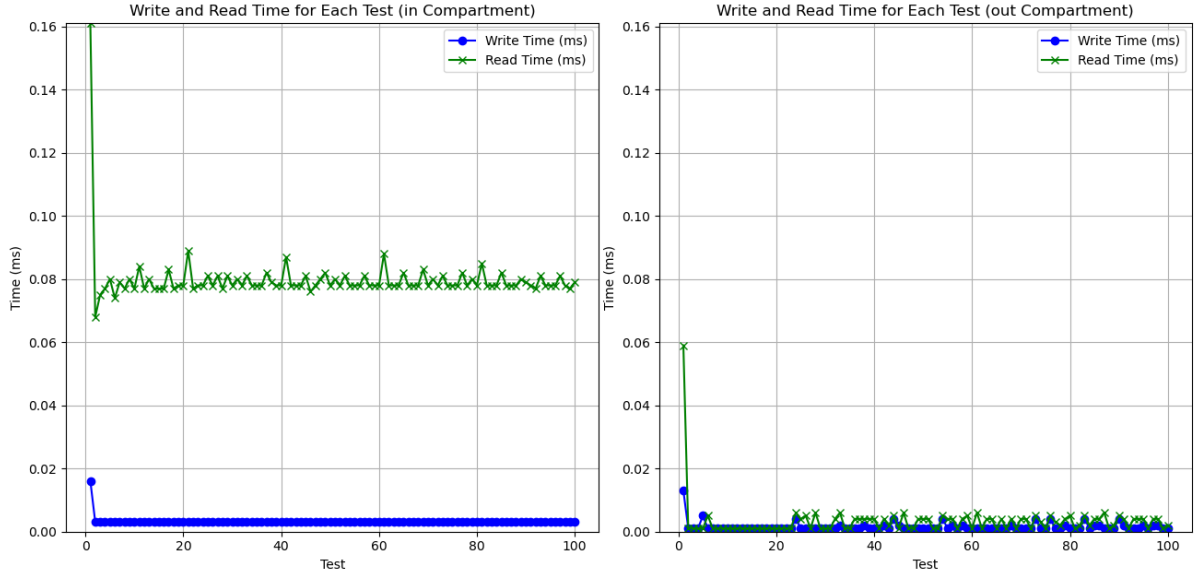
Test	Message Size (Bytes)	Write Time (ms)	Read Time (ms)	Total Time (ms)
1	1024	0.013	0.059	0.072
2	1024	0.001	0.001	0.003
3	1024	0.001	0.001	0.002
4	1024	0.001	0.001	0.002
...	...	...	...	...
100	1024	0.001	0.002	0.003

The Tables 4 and 5 contain information on the test number, the size of the messages, the write and read times, and the total time for each operation.

## Analysing the results

The graphs in the Figure 7 show the write and read times for each test in the two scenarios, inside and outside the secure environment (compartment) on the Morello Board.

As we can see, the graphs reflect variations in write and read times, both inside and outside the safe room environment. There are fluctuations in the times recorded which may point to variability in the pipe's performance under different operating conditions. When comparing the graphs for the two environments, it is possible to notice differences in the times, which suggests that security configurations or differences in the operating environment can significantly influence communication efficiency.



**Figure 7. Comparison of write and read times for pipe communication in secure and normal environments on the Morello Board.**

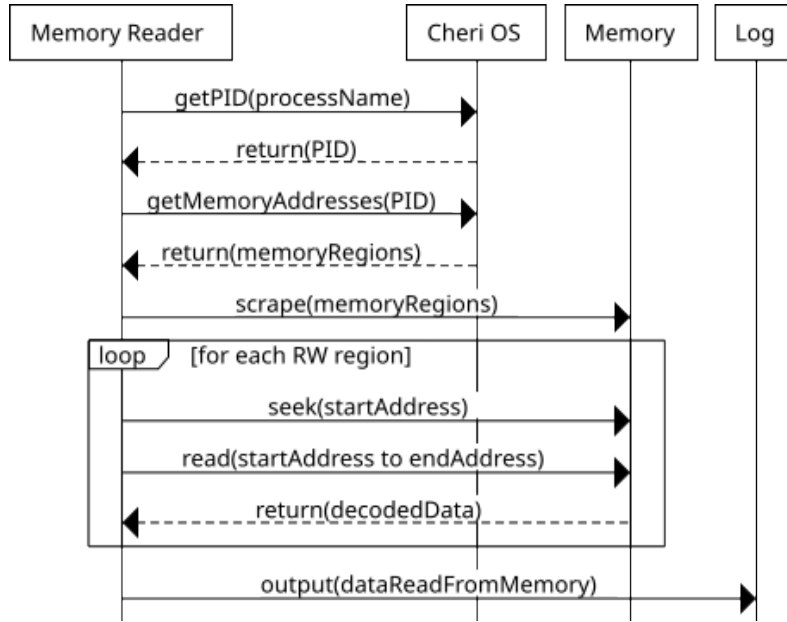
## 5. Testing the Exposure of Sensitive Data in Memory Regions through Direct Memory Access (Memory Scraping)

The main objective of this experiment is to measure and analyze the contents of the memory regions with read and write (RW) permissions of a target program running on the CheriBSD 24.05 operating system. This target program is the same one used in the Attestable’s performance evaluation experiment, `integration_process`. To do this, the experiment identifies the program’s RW memory regions and attempts to directly extract the content stored in these regions, using the Python script `memory_reader.py`, available in the `tee-morello-performance-experiments` repository. This experiment evaluates the potential for exposing sensitive data in memory by directly accessing areas where the program may store temporary or sensitive data.

### Experimental Configuration

The Figure 8 sequence diagram details how a programme’s memory regions can be accessed directly from the operating system. To do this, the process implemented through the Memory Reader code asks the Cheri OS for the PID of the target process by name, `getPID(processName)`, receiving the PID in response. Memory Reader then uses this PID to ask Cheri OS for a list of the memory regions that have read and write (RW) permissions associated with the process, `getMemoryAddresses(PID)`, and the operating system returns the mapped regions. The Memory Reader then starts direct memory reading (scraping) by accessing the process’s memory file and, for each RW region listed, fetches the region’s starting address, reads the content between the starting and ending address, and receives the decoded data. This cycle is repeated for all RW regions, ensuring that all memory areas accessible for reading and writing are explored. Finally, the Memory Reader records the read data in the log, `output(dataReadFromMemory)`.

### Test procedure



**Figure 8. Sequence diagram detailing the memory scraping process through direct access in the Cheri OS.**

To carry out the direct memory reading experiment, Alice performs the following steps on the Morello Board:

1. **Start:** Alice compiles and runs the programme `memory_reader.py` in two different scenarios:
  - **Inside the compartment:** Alice compiles and runs the `integration_process.c` programme on the Morello Board, using the secure environment.
    - **Compile:** `clang-morello -march=morello+c64 -mabi=purecap -g -o integration\_process integration\_process.c -L. -lssl -lcrypto -lpthread`
    - **Run:** `proccontrol -m chericl8n -s enable ./integration\_process`
  - **Outside the compartment:** Alice compiles and runs the `integration_process.c` programme in the Morello Board's normal operating environment.
    - **Compile:** `clang-morello -o integration\_process integration\_process.c -lssl -lcrypto -lpthread`
    - **Run:** `./integration\_process`
2. **Launch:** Alice starts the script that performs direct memory reading with the following command:
  - **Run:** `python3 memory\_reader.py`
3. **Execution:** `memory_reader.py` cycles through each RW region, directly reading the data between the start and end addresses of each region. Alice observes the results on the terminal output.

## Results

The Table 6 shows the results of tests carried out to evaluate access to sensitive data stored and processed by `integration_process` running in different environments - inside and outside Morello's secure enclosure. Each test varied the user's permission level (root or with

reduced permissions), recording whether memory access was successful and whether sensitive data was visible.

**Table 6. Access control test results for sensitive data.**

Test	Environment	User Permissions	Access Result	Sensitive Data Visible
1	in Compartment	Root	Success	Yes
2	in Compartment	Different user	Failure	No
3	out Compartment	Root	Success	Yes
4	out Compartment	Different user	Failure	No

Table 6 metadata:

**Test ID:**

Unique identification for each test carried out.

**Environment:**

Indicates whether the test was carried out inside Morello’s secure compartment or in the normal operating environment.

**User Permissions:**

Details the user’s permission level (e.g., root or a user with reduced permissions).

**Access Result:**

Result of the access, indicating whether the memory region was successfully accessed.

**Sensitive Data Visible:**

Indicates whether the test was able to visualize sensitive data or not.

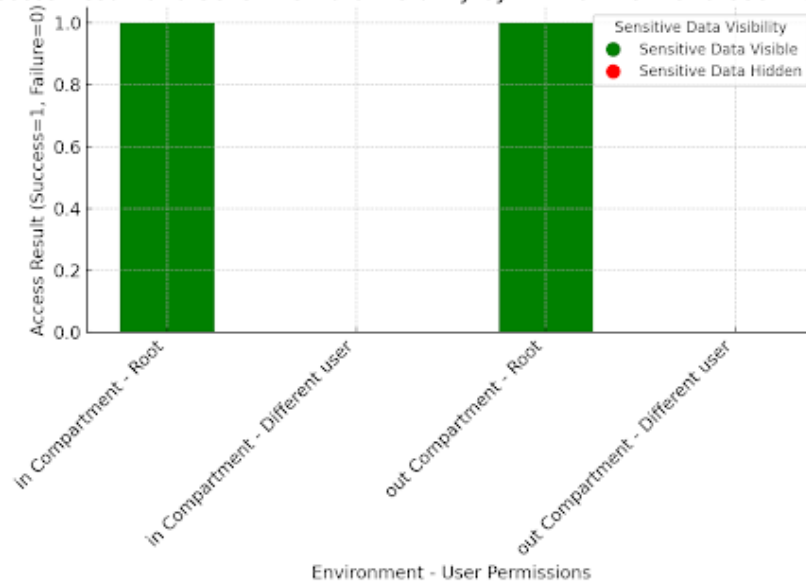
**Analysing the results**

The graph in the Figure 9 shows the differences in the results of access and visibility of confidential data according to the environment and user permissions.

The results show that access to sensitive data is allowed to any user with root permissions, both in Morello’s secure compartment and outside it, indicating that high permissions grant unrestricted access. In contrast, users with reduced permissions are unable to access this data in either environment, demonstrating that the secure enclosure consistently blocks unauthorised access, while the environment outside the enclosure also maintains effective controls.

However, although compartmentalisation is designed to protect sensitive data by restricting access, including to the user who initiated the process, the results indicate that the current configuration may not be sufficient to guarantee complete isolation. This behaviour highlights the importance of investigating further adjustments to the Morello Board’s compartmentalisation, especially for contexts where there is a need to protect sensitive data.

Access Result and Sensitive Data Visibility by Environment and User Permissions



**Figure 9. Differences in access results and visibility of confidential data based on environment and user permissions.**



## 6. PRINTS

Process memory regions running outside the secure compartment:

PID	START	END	PROT	RES	PRIS	RSP	SHD	FLAG	TP	PATH
3701	0x200000	0x202000	r--R-	2	0	3	1	0----	vn	/home/regis/tee-compartmentalisation-study-case/launcher/programs-data-base/sources/out_server
3701	0x211000	0x214000	r-xR-	3	0	3	1	0----	vn	/home/regis/tee-compartmentalisation-study-case/launcher/programs-data-base/sources/out_server
3701	0x223000	0x224000	r--R-	1	0	1	0	C-----	vn	/home/regis/tee-compartmentalisation-study-case/launcher/programs-data-base/sources/out_server
3701	0x233000	0x234000	rw-RW	1	1	1	0	-----	sw	
3701	0x88a19000	0x88a29000	-----	0	0	0	0	G-----	pd	
3701	0x88a29000	0x88a39000	rw-RW	5	5	1	0	---B-C	sw	
3701	0x1071000	0x1072000	r-x---	1	1	42	0	-----	ph	
3701	0x1cc58000	0x1cc3000	r--R-	51	168	18	4	0----	vn	/usr/lib64/libssl.so.30
3701	0x1cc3000	0x1cc2000	-----	0	0	0	0	0----	pd	
3701	0x1cc2000	0x1cc1000	r-xR-	85	168	18	4	0----	vn	/usr/lib64/libssl.so.30
3701	0x1cd17000	0x1cd16000	-----	0	0	0	0	0----	pd	
3701	0x1cd16000	0x1cd15000	r--R-	9	0	1	0	C-----	vn	/usr/lib64/libssl.so.30
3701	0x1cd15000	0x1cd14000	-----	0	0	0	0	0----	pd	
3701	0x1cd14000	0x1cd13000	rw-RW	6	0	1	0	C-----	vn	/usr/lib64/libssl.so.30
3701	0x12064000	0x12072000	r--R-	13	42	18	4	0----	vn	/usr/lib64/libthr.so.3
3701	0x12071000	0x12060000	-----	0	0	0	0	0----	pd	
3701	0x12060000	0x12059000	r-xR-	22	42	18	4	0----	vn	/usr/lib64/libthr.so.3
3701	0x12059000	0x12058000	-----	0	0	0	0	0----	pd	
3701	0x12058000	0x12057000	r--R-	1	0	1	0	C-----	vn	/usr/lib64/libthr.so.3
3701	0x12057000	0x12056000	-----	0	0	0	0	0----	pd	
3701	0x12056000	0x12055000	rw-RW	2	0	1	0	C-----	vn	/usr/lib64/libthr.so.3
3701	0x12055000	0x12054000	rw-RW	10	38	1	0	-----	sw	
3701	0x12054000	0x12053000	r--R-	93	392	18	4	0----	vn	/usr/lib64/libc.so.7
3701	0x12053000	0x12052000	-----	0	0	0	0	0----	pd	
3701	0x12052000	0x12051000	r-xR-	268	392	18	4	0----	vn	/usr/lib64/libc.so.7
3701	0x12051000	0x12050000	-----	0	0	0	0	0----	pd	
3701	0x12050000	0x12049000	r--R-	10	0	1	0	C-----	vn	/usr/lib64/libc.so.7
3701	0x12049000	0x12048000	-----	0	0	0	0	0----	pd	
3701	0x12048000	0x12047000	rw-RW	7	0	1	0	C-----	vn	/usr/lib64/libc.so.7
3701	0x12047000	0x12046000	rw-RW	17	17	1	0	-----	sw	
3701	0x13f76000	0x13f75000	r--R-	389	648	18	4	0----	vn	/usr/lib64/libcrypto.so.30
3701	0x13f75000	0x13f74000	-----	0	0	0	0	0----	pd	
3701	0x13f74000	0x13f73000	r-xR-	431	648	18	4	0----	vn	/usr/lib64/libcrypto.so.30
3701	0x13f73000	0x13f72000	-----	0	0	0	0	0----	pd	
3701	0x13f72000	0x13f71000	r--R-	82	0	1	0	C-----	vn	/usr/lib64/libcrypto.so.30
3701	0x13f71000	0x13f70000	-----	0	0	0	0	0----	pd	
3701	0x13f70000	0x13f69000	rw-RW	8	0	1	0	C-----	vn	/usr/lib64/libcrypto.so.30
3701	0x13f69000	0x13f68000	rw-RW	1	1	1	0	-----	sw	
3701	0x13f68000	0x13f67000	-----	0	0	0	0	0----	pd	
3701	0x13f67000	0x13f66000	rw-RW	31	31	1	0	-----	sw	
3701	0x13f66000	0x13f65000	rw-RW	52	52	1	0	-----	sw	
3701	0x13f65000	0x13f64000	rw-RW	31	31	1	0	-----	sw	
3701	0xffffffff	0x300000000000	-----	0	0	0	0	G-----	pd	

Figure 10. Process memory regions running outside the secure compartment.

### Memory regions of the process running inside the secure compartment:

[illegible]

**Figure 11. Memory regions of the process running inside the secure compartment.**

## Finalisation of the process by the system

The process runs for a certain period of time (around 1 hour) and is finalised by the system: killed.

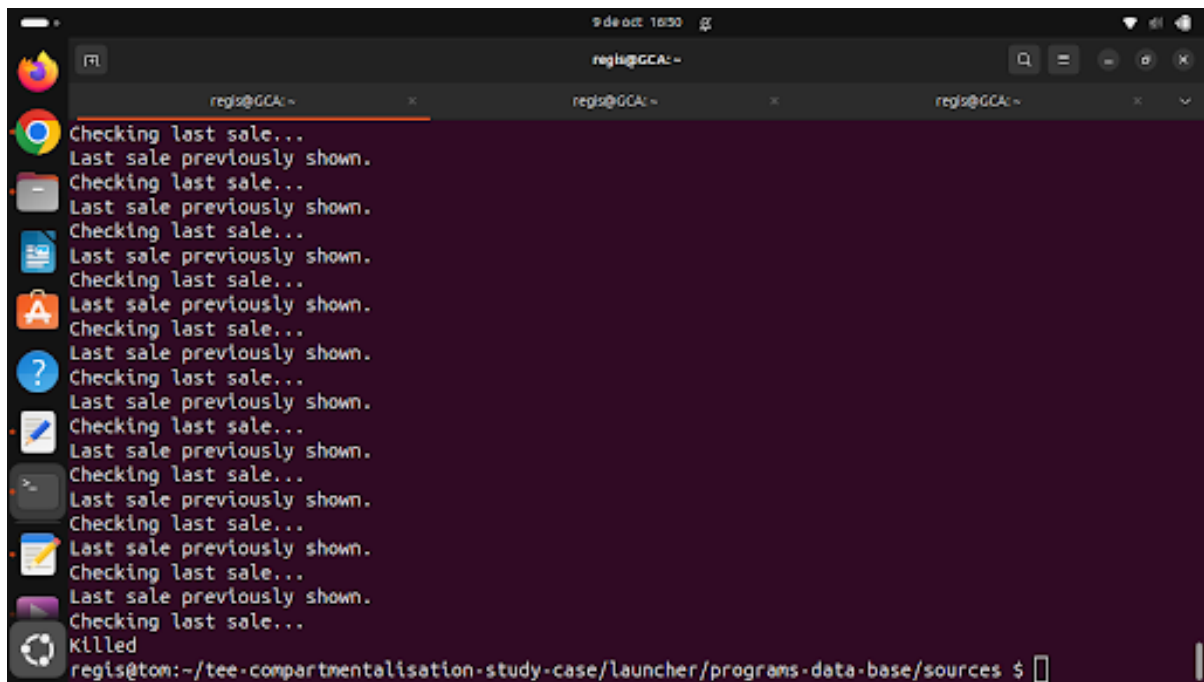


Figure 12. Finalisation of the process by the system.

**The morello board crashed at this point.**

Regis needs to contact the morello-board s administrator [erik] to reboot.

3587	0x43000000	0x433e0000	-----	0	0	0	0	G-----	gd
3587	0x433e0000	0x43400000	rw-RW	3	3	1	0	---D-c	sw
3587	0x43400000	0x43680000	rw-RW	583	101679	21	0	-----c	sw
3587	0x43680000	0x43980000	rw-RW	764	101679	21	0	-----c	sw
3587	0x43980000	0x43d00000	rw-RW	882	101679	21	0	-----c	sw
3587	0x43d00000	0x44100000	rw-RW	997	101679	21	0	-----c	sw
3587	0x44100000	0x44600000	rw-RW	1277	101679	21	0	-----c	sw
3587	0x44600000	0x44c00000	rw-RW	1485	101679	21	0	-----c	sw
3587	0x44c00000	0x45300000	rw-RW	1788	101679	21	0	-----c	sw
3587	0x45300000	0x45b00000	rw-RW	2044	101679	21	0	-----c	sw
3587	0x45b00000	0x46500000	rw-RW	2559	101679	21	0	-----c	sw
3587	0x46500000	0x47100000	rw-RW	3065	101679	21	0	-----c	sw
3587	0x47100000	0x47f00000	rw-RW	3584	101679	21	0	-----c	sw
3587	0x47f00000	0x48f00000	rw-RW	4074	101679	21	0	-----c	sw
3587	0x48f00000	0x4a300000	rw-RW	5114	101679	21	0	-----c	sw
3587	0x4a300000	0x4bb00000	rw-RW	6120	101679	21	0	-----c	sw
3587	0x4bb00000	0x4d700000	rw-RW	7163	101679	21	0	-----c	sw
3587	0x4d700000	0x4f700000	rw-RW	8124	101679	21	0	-----c	sw
3587	0x4f700000	0x51f00000	rw-RW	10225	101679	21	0	-----c	sw
3587	0x51f00000	0x54f00000	rw-RW	12119	101679	21	0	-----c	sw
3587	0x54f00000	0x58700000	rw-RW	14025	101679	21	0	-----c	sw
3587	0x58700000	0x5c700000	rw-RW	14359	101679	21	0	-----c	sw
3587	0x5c700000	0x61700000	rw-RW	1328	101679	21	0	-----c	sw
3587	0xfbfdbffff000	0xfbfdc0000000	rw---	1	0	1	0	C----c	sw
3587	0xfbfdc0000000	0xfe0000000000	rw---	848	0	1	0	C-----	sw
3587	0xffffbfeff000	0xffffbfff80000	rw-RW	1	1	1	0	CN---c	sw
3587	0xffffbfff80000	0xfffffff60000	-----	0	0	0	0	G-----	gd
3587	0xfffffff60000	0xfffffff80000	rw-RW	3	3	1	0	C--D-c	sw
3587	0xfffffff80000	0x1000000000000	r-x--	1	1	44	0	-----	ph

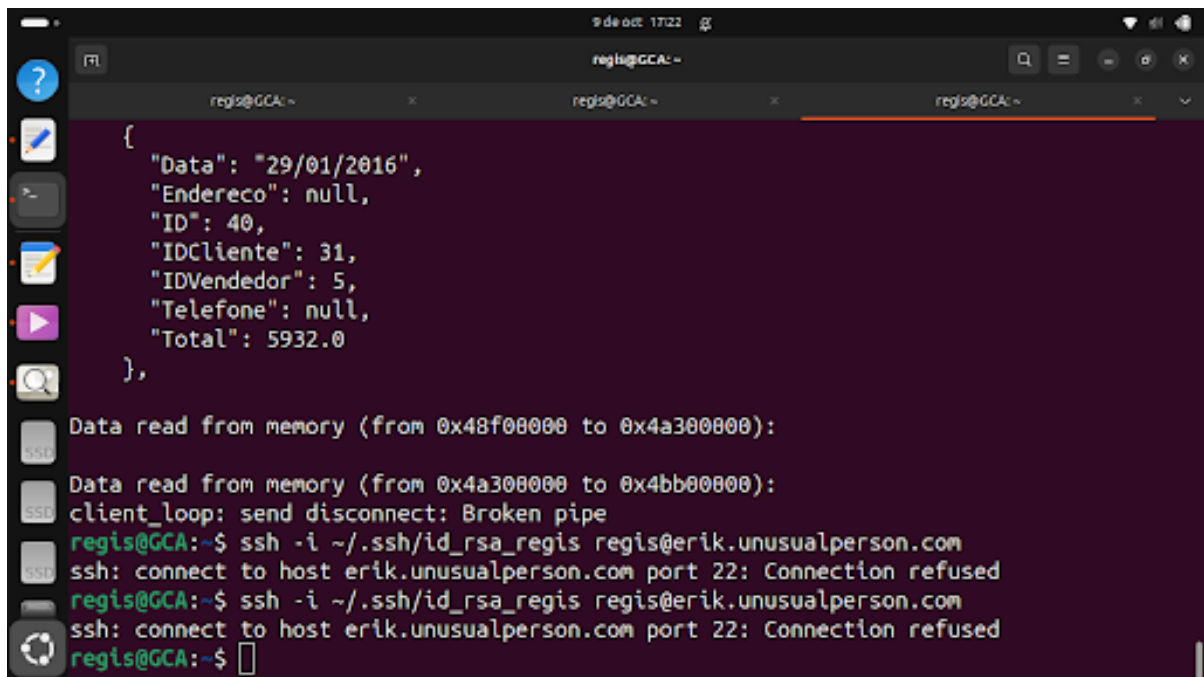
**Figure 13. The morello board crashed at this point.**

The remote ssh shell that connects to the morello board crashes when the mem scanner prog tries to read this range of mem add: 0x4a300000 - 0x4bb00000.

Other ssh shell connections continue working as normal.

There are occasion when the prog manage to crash the whole morello board [the actual cheriBSD] , no more shells can be opened. It seem that total crash takes place when the prog tries to read some specific mem ranges [we dont know which ones cause total crashes; we believe that they are ranges where priviledge soft runs.

This is our preliminary observation.



```
regis@GCA: ~  
{  
  "Data": "29/01/2016",  
  "Endereco": null,  
  "ID": 40,  
  "IDCliente": 31,  
  "IDVendedor": 5,  
  "Telefone": null,  
  "Total": 5932.0  
},  
Data read from memory (from 0x48f08000 to 0x4a308000):  
Data read from memory (from 0x4a308000 to 0x4bb08000):  
client_loop: send disconnect: Broken pipe  
regis@GCA:~$ ssh -i ~/.ssh/id_rsa_regis regis@erik.unusualperson.com  
ssh: connect to host erik.unusualperson.com port 22: Connection refused  
regis@GCA:~$ ssh -i ~/.ssh/id_rsa_regis regis@erik.unusualperson.com  
ssh: connect to host erik.unusualperson.com port 22: Connection refused  
regis@GCA:~$
```

Figure 14. client\_loop: send disconnect: Broken pipe.



## Error occurs after server crash and restart [Erik].

### Error accessing process memory:

```
[Errno 2] No such file or directory: '/proc/3587/mem'
```

This error occurs because the file `/proc/{pid}/mem`, which the script tries to access to read a process's memory, is not available or cannot be accessed. This may happen if the process does not exist, the PID is incorrect, or the script does not have the necessary permissions to access this path.

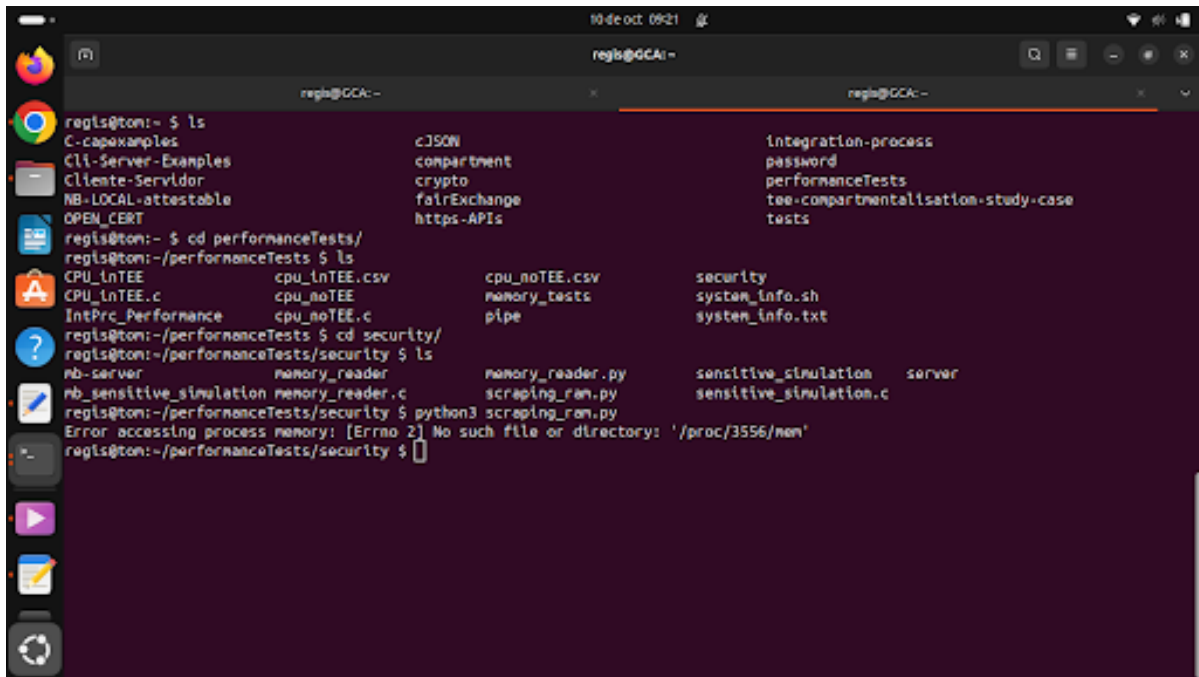


Figure 15. [Errno 2] No such file or directory: '/proc/3587/mem'.

## Procedure for running `memory_reader.py` after MB is rebooted.

Check that `/proc` is mounted correctly:

```
regis@tom:~/performanceTests/security $ mount | grep /proc
```

Figure 16. Check that `/proc` is mounted correctly.

If it's not mounted, you need to try mounting it:

```
regis@tom:~/performanceTests/security $ sudo mount -t procfs proc /proc
```

Figure 17. If it's not mounted, you need to try mounting it.

After the command to mount `/proc`, it was simply possible to run the `memory_reader.py` script again.